

Visualizing Geo-tagged Data

Ján Vorčák *janvor@ifi.uio.no*, Frode Tobias Bjerke *frodetbj@ifi.uio.no*,
Matthäus Skiba *mskiba@mail.uni-mannheim.de*

Abstract—The thirst for knowledge nowadays motivates people to search the internet for different systems to satisfy their demand. Google Maps, Google Street View or Flickr are used to get impressions of a specific region whereas Wikipedia can deliver the appropriate facts. However, the user has to utilize at least two systems. Furthermore, he has to search for additional information on his own.

In this paper, we describe how we extended a system that creates interactive video tours by combining videos with location data. We show how relevant nodes can be identified and how additional information about these points can be retrieved. Furthermore, we describe how this information can be displayed by the system.

Keywords—Geotagging, Interactive Video Tour, Video Navigation.

I. INTRODUCTION

Almost all of us have experienced this situation: You want to visit a city or you have an appointment in a part of a city you are not familiar with. It becomes more and more popular to have a look at Google Maps to plan your route or to check up the surroundings at Google Street View. Even tagged photos on Flickr will be consulted to improve the estimation of the environment. Furthermore, it happens that the user wants more specific information about points of his interest. For sure, he can search for further information on his own but this is still related to some effort. Imagine the user sees a statue on Google Street View in a city he wants to visit. How should he get further information about this statue, since he has to know at least a correct spelled name? He also might be interested in places or points within his surrounding, although he cannot see them from his actual position or as he even might be unaware about their existence.

It is desirable to have a system that can fulfill the above needs. It should show routes, the environment as well as interesting points in the surrounded area. Information about these points shall also be included. Therefore, the system has to be aware of the current position, so it can query for additional points and information.

Since we try to implement such a system, we will build up on a system, which creates interactive video tours. The described system combines videos with location data. A nice side-effect is that videos carry more information than pictures. Furthermore, it provides the current position since it is necessary to navigate through the video or rather through the path. With this information we are able to query available databases for nearby points and their entities. Since such databases exist, it is not necessary to create one by ourselves. This would be a not manageable task.

The contribution of this work is to show how points of interest can be found towards the current position and how

information according to these points can be retrieved and visualized.

This report is structured as follows. Related works shows some work that is done on geo-tagged multimedia, geo-tagged multimedia also gets classified. Design we talk about how we envision our solution. Implementation we talk about how we implemented our design and comment on some issues we had. In the evaluation chapter we further comment on aspects of our solution. The evaluation will only comment on the features we have added to the existing system. Last, we conclude our work in the conclusion chapter.

II. RELATED WORK

Luo et al[1] propose to split geotagged multimedia into four different modalities: points of interest geographical databases, aerial image databases, unstructured geo-referenced web resources and collection of geo-referenced multimedia.

Sources serving data from the points of interest geographical databases modality are often of high reliability and provides data on things like landmarks, shops and other points of interest. Interesting sources for such data could be OpenStreetMap Overpass API[2] and Geonames[3].

Aerial images covering the earth has been knit together into great interactive services such as Google Earth[4]. This kind of imagery can cover large areas and are good at visualizing larger areas of land.

Unstructured geo-referenced web resources could come from websites, news, blogs, wikis or social networks. This is a very diverse modality where data could belong to any number of different modalities. A great strength of this modality however how diverse is presenting live information on events or major happenings. Data is often provided by non-experts and could come in tremendous volumes, this is certainly the fact for social networks. Hence, data in this modality could be noisy and lack of objectivity. Therefore, heavy pre-processing is often required to present reasonable results. Some major resources here are Twitter[5] and Facebook[6].

Treating geo-referenced multimedia as collections as opposed to single pieces makes it possible to find more relevant information by means of aggregation. There have specially been done a lot of aggregation on geo-referenced images. Images on large photo-sharing services include metadata such as tags, likes and comments; data mining techniques can therefore be applied to aggregate interesting information on locations and areas.

Spyrou and Mylonas[7], and Kisilevich et al[8] worked with the geo-referenced multimedia as collection modality. A image source they used was for example Flickr[9]. Spyrou and Mylonas' work focused on finding the most relevant tags for an area represented by a cluster of images. Tags

are usually user-created and therefore of varying quality, in their work they proposed a tag-ranking model that would find the most descriptive tags for a given area. Kisilevich et al's work on the other hand focused on locating and visualizing what they referred to as attractive places. They located areas by image-clusters before applying methods of analyzing their attractiveness. Among their techniques to find the best places they considered the popularity of the pictures in a cluster.

Chippendale et al[10] worked cross modalities, what they did was showing ways to explore and organize photos. Two of the more interesting things they present in their work are augmenting photos with any other form of geo-tagged data and photo alignment. In augmenting photos they added markers as form of an icon representing the source. Sources could be anything geo-tagged such as news, other photos or videos. In one example they had a picture overlooking a valley which had been augmented with YouTube icons on different locations in the valley. Each Icon indicated a position where a YouTube video was geo-tagged. To help show the position of the YouTube video in the image, the icons representing each video had different sizes according to the distance from where original image was taken.

III. DESIGN

Our task is to enrich an already existing system[11] with geo-tagged multimedia. The existing system, as seen in figure 1, visualizes geo-tagged video according to position rather than time. It works by users uploading geo-tagged video. By geo-tagged in this specific context it means that the video is accompanied with location data, the location data tells which location the camera is at throughout the video. In the web application, videos can then be navigated by a map where the tracks of each video is visualized.

We have two main goals we want to achieve with our design. First, visualize points of interest on the map. There should be ways to find relevant information on a point of interest in the application. Secondly, the points of interest should show up on the videos. Such that when a video-clip passes a certain building which is marked as a point of interest a label indicating what point of interest the given building is.

Regarding finding points of interest this will be done in two steps. First, we have to think about points we consider as relevant and that we want to receive. In the following, we will call these points of interest for POI. The second step is to get additional information from different data sources referring to the POI. As a result, it is necessary to define an interface which provides additional information relating to the POI to the application. Ideally, the interface is able to request this information from different sources in different data and media types.

In the video, the same POIs as in the map should be visualized when they are in close proximity in front of the camera. The position of a POI should be indicated by if it is shown on the left or the right side of the video. Additionally the distance to the POI can be shown by changing the size of the indicator.

IV. IMPLEMENTATION

A. Finding points of interest

The given existing application uses OpenStreetMap (OSM) to visualize available tracks and the current position on the video. It suggests itself to use the available OSM API but according to its documentation, the API is dedicated for editing and not for downloading [12]. Mass requests tend to be resource intensive and since we want to respect this appeal, we use the suggested Overpass API to query for our POI.

The Overpass API supports different Query languages like an XML query format and an own, more concise, query language [13]. We choose to build our queries in the XML format since it is more human readable. The resulting query can be transformed into a concise Overpass query using the given converter. The benefit of using the Overpass query is that it can be performed through a single HTTP request to the interpreter. However, one should pay attention to the number of semicolons at the end of the query as we discovered a little bug due to this work.

Since OSM represents everything with nodes and each node has to be defined by a key-value-pair [14], we query for our POI with these key-value-pairs according to a certain radius to our current position. The current position is provided by the system with the method call and is needed to build a little bounding box to get the node of the current position from Overpass. Since the box is formed by only two coordinates, it represents the actual point we are looking for. Listing 1 and Listing 2 show the used query in the current project in XML format as well as in Overpass Query Language.

Listing 1. Overpass query in XML format

```
1 <osm-script output="json">
2   <bbox-query s="49.489696" n="49.489696" w
   = "8.467391" e="8.467391"/>
3   <query type="node">
4     <around radius="1000"/>
5     <has-kv k="historic"/>
6   </query>
7   <print/>
8 </osm-script>
```

Listing 2. Overpass query in Overpass Query Language

```
1 [out:json];node
  (49.489827,8.468115,49.489927,8.468215);node
  (around:1000)["historic"];out;
```

The query from Listing 1 retrieves nodes with the key "historic" and every belonging value (e.g.: monuments, memorials, ruins or shrines) that are around a thousand meters from the Marketplace (G1) in Mannheim. The result will be delivered as a JSON object which will be passed on to other methods in the system that take care of displaying or getting further information. This query represents just a proof of concept but can be easily extended by appending further <query>-tags like in Listing 3. This query will have the same result as it would be concatenated with an OR-operator. Adding more elements, which can be considered as arguments in this case, into the same <query>-tag will have the same result as an AND-operator.

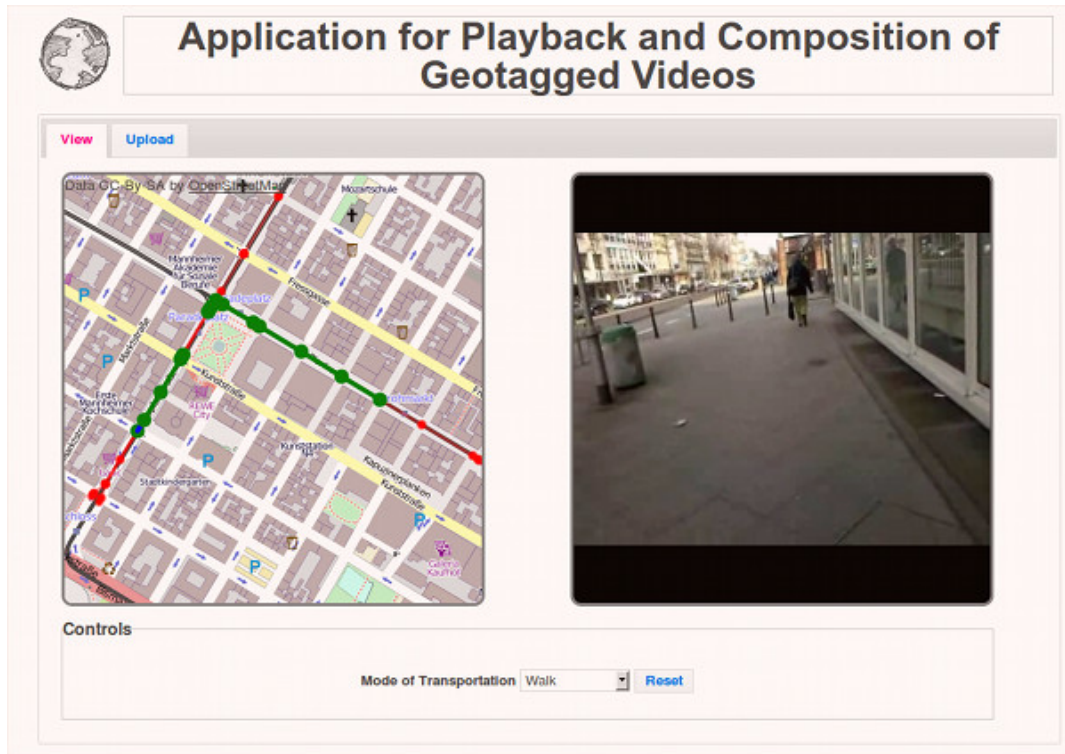


Fig. 1. Screenshot of original application. Retrieved from [11].

The Overpass API offers much more possibilities like membership relations between nodes, regular expressions within key-value pairs and unions between queries. Since we consider the explained functions as sufficient to get all interesting geographical points near a given position, we will not go into detail with these sophisticated functions as this is not an introduction to the Overpass API.

Listing 3. Enhanced Overpass query in XML format

```
1 <osm-script output="json">
2   <bbox-query s="49.489696" n="49.489696" w
   = "8.467391" e="8.467391"/>
3   <query type="node">
4     <around radius="1000"/>
5     <has-kv k="historic"/>
6   </query>
7   <print/>
8   <query type="node">
9     <around radius="1000"/>
10    <has-kv k="leisure" v="park"/>
11  </query>
12  <print/>
13  <query type="node">
14    <around radius="1000"/>
15    <has-kv k="amenity" v="theatre"/>
16  </query>
17  <print/>
18 </osm-script>
```

B. Add data to points of interest

After we download information about the objects nearby, we would like to know more about them and display more

information from external sources like Wikipedia or Flickr.

Our aim is to provide an interface capable of doing request for more external APIs. Interface will be implemented in Javascript and will use server support if needed.

Javascript interface is illustrated in Listing 4.

Listing 4. Javascript interface for connecting external API

```
1 $.info.get = function (api_name, title ,
   parameters)
```

So in order to get Wikipedia content for object with name *NameOfTheObject* you need to call code in Listing 5.

Listing 5. Getting Wikipedia content

```
1 $.info.get("wikipedia", "NameOfTheObject")
```

In order to call Wikipedia API, we wanted to contact its API directly from the Javascript code. The first problem is that we are requesting data from different domain, so in order to avoid cross-domain problems we used *JSONP* to get the data.

After solving this issue, data we are getting are still in the Wikimedia format, which needs to be parsed by the client.

At the end, we decided to have a support for the Wikipedia API on the server side, because we can use Python packages like *wikitoools* to work with Wikipedia API in a more elegant way.

Data we gather from the APIs will be part of the tooltips visible in the video and on the map. For now we support Wikipedia API, but it's possible to easily add other APIs.

Each API will have to be treated individually, some of them will require server support also because of the API keys, some

of them will require caching of the data because of the API limitations. What we require from the future implementations is to follow the defined interface, so that developers can always rely on that kind of abstraction.

C. Displaying POIs on the map

This part drafts how the system displays points of interest on the map. It was not sufficient enough to implement a method that adds markers to the map as changes to the core of the system were needed. First of all, it was necessary to add another layer for the markers that should be drawn by the OpenLayers API. Furthermore the depth of this layer had to be set, so a popup would appear after a click on the marker. If the Layer is not correctly initialized, the application will not recognize any motion on this layer.

As these changes were performed, the marker drawing function *setMarker* could be implemented. This function is called from the success function of the performed AJAX call in *onVideoProgress* after the POIs were received from the Overpass API. The points the system is interested in are stored in a JSON object. This object will be iterated to get the longitude, the latitude and the name. These are the parameters that are required to create and set a marker with an associated popup. Afterwards the method adds an event to the popup in case the mouse moves over the marker.

All in all, after the core of the system has been modified, markers can be added to the map very easily. However, the documentation of the OpenLayers API is very poor. This fact turned this easy looking part of displaying markers on a map into a real challenge.

D. Displaying POIs in the video

Once we have all the data available locally in the web browser, it's time to visualize them on the video.

In the existing system, the HTML5 video tag is used for displaying the video content. The idea is to add HTML overlay over this video and display custom content there. This task thus consists of three different sub-tasks.

- Overlaying HTML5 video tag
- Filtering the places which are visible in the video
- Displaying data in the video overlay

1) *Overlaying HTML5 video tag*: For this task we will just construct *div* tag with id *video-wrapper* with *relative* positioning. This wrapper will contain two elements - video tag and another *div* tag with id *video-overlay*.

Since video wrapper will have relative positioning, all we need is to make sure that video overlay has higher *z-index* and we can add objects here which will be visible above the video tag.

The only complication which can happen with this solution is bad behavior during fullscreen mode.

If we want to make video fullscreen, we would loose the objects, because they are just placed in another layer. Solution is to use HTML5 *requestFullscreen* function. If we want to make video fullscreen, we will need to call HTML5 *requestFullscreen* function on the video wrapper element, since

it contains both the video and the overlay. It's important to mention that *requestFullscreen* function is only triggerable via user input (i.e. clicking the button), because of the security.

2) *Filtering visible places*: We have already fetched all the interesting information about objects nearby. If we want to now display them above the video, we need to know which objects are actually nearby or visible.

If we know which places are nearby, we can display some tips in the video. For the objects that can be visible on the video, we will try to estimate their location in the video.

In the objects fetched from the external API, we are provided with the coordinates of each object. Since we need to recalculate the current view every time we have moved in the video, we will create a function *showPlacesOnTheVideo* which will be called from within already provided *onVideoProgress* function.

This function is great place to do these kind of calculations, because it is called every time the video proceeds a new point on the map. Updating the view on the video in this interval is thus sufficient. If more precise calculations will be needed in the future, it is possible just to call *showPlacesOnTheVideo* function more times using the timer.

In the *onVideoProgress* function, we are already provided with the current point and the next point we are about to reach. Now we know where are we located, our direction and places nearby.

a) *Getting the places nearby*: In order to get all the nearby places, we need to iterate through them and check if

$$abs(current.lon - place.lon) \leq r$$

and

$$abs(current.lat - place.lat) \leq r$$

where *current* is object containing our coordinates, *place* is the object containing the coordinates of the place we are interested in and *r* is a constant defining the acceptable difference for the coordinates, which is set to 0.0009 by default.

b) *Getting the places in front of us*: As seen on Figure 2, we will know the coordinates of the current point and the point we are about to reach. What we want to count are points A, B, C and D. If we can calculate the positions of these points, we will be able to say that the object is in front of us, if it's located within this rectangle.

We will be also able to say whether an object is located on the right side of the video or the left side of the video, depending on which of the areas of the rectangle contains the object. We will be also able to return the distance from the object, which can be useful while rendering the object on the video.

After recalculating, we will be able to say whether the object is located on the right or left side of the street and say how far away it is. Which should be sufficient for estimating its position within video overlay.

In order to do this we need to first calculate vector of our direction vector.

$$CURRENT = [lat, lon]$$




Fig. 2. Calculating the places in front of us

$$NEXT = [lat, lon]$$

$$\vec{direction} = (CURRENT, NEXT)$$

When we have the direction vector, we can calculate the normal vectors.

$$\vec{n1} = (-direction[1], direction[0])$$

$$\vec{n2} = (direction[1], -direction[0])$$

If we add normal vectors to the points *NEXT* and *CURRENT* we can calculate the coordinates of *A*, *B*, *C* and *D*

$$A[0] = NEXT[0] + n1[0] * c$$

$$A[1] = NEXT[1] + n1[1] * c$$

where *c* is the constant specifying the width of our view.

Once we have calculated all of the coordinates, we can say whether object we are trying to locate is within the range and if so we can analyze whether it is on the right side or left side of the view.

Since we know the *CURRENT* coordinates and coordinates of the video we can also return the distance of the objects from our current location.

In case we want to specify wider or longer view range. We can move the *NEXT* point by adding the direction vector for longer range or we can specify higher *c* constant for wider range.

Constant *c* is initially set to 1, which means that distance between points *CURRENT* and *NEXT* is the same as distance between points *NEXT* and *B*.

3) *Displaying data in the video overlay*: After we are done with calculations and we can say about each object whether it is located in front of us or not, we can start visualizing these elements on the video overlay.

Properties we have determined for each object are it's distance, information whether it is on the right or left side of the street and other properties like name and type of the object we fetched from the external APIs.

It is important to realize that video may have different dimensions so all of the positioning needs to be done in the relative way.

We will set CSS *float* attribute to *left* or *right* for all the objects depending on their positions and will set up their size and position depending on the distance of the object.

Our assumption is that this system will be used mostly in streets so specifying that something is on the right side or the left side should be sufficient.

V. EVALUATION

A. Historical locations marked on the map

For our evaluation our POI finder was only configured to find historical locations. Every POI found inside a preset radius of our location was displayed neatly on the map with markers. Intuitively, as seen from other interactive map services like Google Maps [15], clicking the marker would show whatever information we have gathered about the location in a popup.

A problem we encountered however was that at some locations there could be several POIs in very close proximity or on each other. Solving how to visualize all the different POIs is something that should be examined in further work.

The information we gathered about a location for this evaluation was only the Wikipedia entry found by querying the name of the POI. As the system was configured to only show one popup at the time, the screen never became cluttered with popups. A popup could easily be closed by a close button in the top right corner, indicated by a red x-icon.

B. Language issues on historical locations

We tested our system with videos from Mannheim in Germany. When querying for historical points of interest from Overpass[2] our result set consisted of only german names for places. When displaying information in popups on the map we combine the name from points of interest with text retrieved from Wikipedia API[16]. Wikipedia was queried by using the name of the point of interest. Figure 3 show a popup where english Wikipedia content was found. As the names of the points of interest where in german and we query the english Wikipedia we rarely got any hits. Figure 3 was actually one of the very few points of interest with english wikipedia content. Figure 4 displays a popup where there was not found any Wikipedia information.

C. The POI indicators on videos

For the evaluation of displaying content above the video, we used the historical locations fetched from OpenStreetMap Overpass API [2].

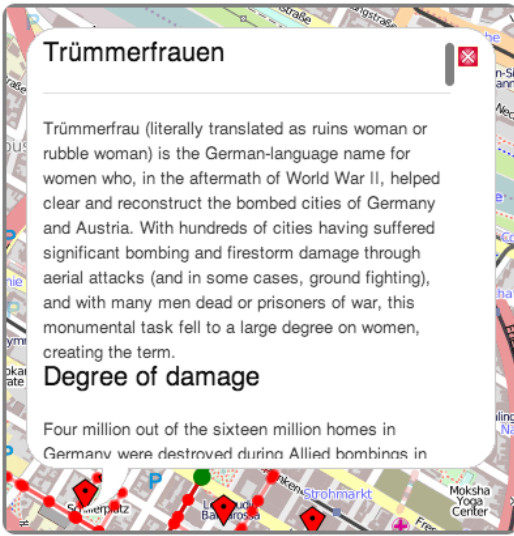


Fig. 3. Point of interest popup with english Wikipedia content.



Fig. 4. Point of interest popup where no english Wikipedia content was found.

In the scenario visible in Figure 5, there are several routes where we could find historical locations on our way. We choosed all the possible routes from this scenario and we found out that we are notified with the relevant alert above the video layer in case interesting venue is nearby.

In addition, the location of the venue is estimated using the hint whether the object is located on the right or the left side of the street.

The hints are added and removed from the video dynamically, so we are notified about the objects currently visible in the video.

Figure 5 illustrates an example of the interesting historical venue near our current location.

In case of more venues are visible on the video, we are




Fig. 5. Example of the content displayed in the HTML5 video

notified with several alerts above the video layer on both sides.

VI. CONCLUSION

We have in this work enriched an existing interactive geo-location video tour application.

We find exiting places in an area before we find additional information on this location. Our implementation for finding additional information is made in such a way that extending the system with further data sources should be tricial.

Information and location of the interesting locations can be viewed from the map. Additionally, points of interest are marked in the geo-tagged videos as well. When a video-tour passes a point of interest, this is indicated in the video.

This application really reveals the power of geo-tagged information; as wikipedia information seamlessly is intertwined with geo-tagged videos into one simple to use interactive application. As the system is easy extendable for more sources of information, one can only start imagining how this could be extended with further imagery, YouTube[17] videos and live social media content.

REFERENCES

- [1] Luo, J., Joshi, D., Yu, J., and Gallagher, A. (2011). Geotagging in multimedia and computer vision? a survey. *Multimedia Tools and Applications*, 51(1), 187-211.
- [2] (2013, Apr.) OpenStreetMap Overpass API. http://wiki.openstreetmap.org/wiki/Overpass_API
- [3] (2013, Apr.) Geonames. <http://www.geonames.org/>
- [4] (2013, Apr.) Google Earth. <http://www.google.com/earth/index.html>
- [5] (2013, Apr.) Twitter. <http://www.twitter.com>
- [6] (2013, Apr.) Facebook. <http://www.facebook.com>
- [7] Spyrou, E., and Mylonas, P. (2011, December). Placing User-Generated Photo Metadata on a Map. In *Semantic Media Adaptation and Personalization (SMAP)*, 2011 Sixth International Workshop on (pp. 79-84). IEEE.

- [8] Kisilevich, S., Mansmann, F., Bak, P., Keim, D., and Tchaikin, A. (2010, February). Where would you go on your next vacation? A framework for visual exploration of attractive places. In *Advanced Geographic Information Systems, Applications, and Services (GEOPROCESSING)*, 2010 Second International Conference on (pp. 21-26). IEEE.
- [9] (2013, Apr.) Flickr. <http://www.flickr.com/>
- [10] Chippendale, P., Zanin, M., and Andreatta, C. (2009, November). Collective photography. In *Visual Media Production, 2009. CVMP'09. Conference for* (pp. 188-194). IEEE.
- [11] (2013, Apr.) <http://ls.wim.uni-mannheim.de/de/pi4/research/projects/geotagged-videos>
- [12] (2013, Apr.) OpenStreetMap wiki about how to download data. [Online]. Available: http://wiki.openstreetmap.org/wiki/Getting_Data
- [13] (2013, Apr.) Overpass API language guide. [Online]. Available: http://wiki.openstreetmap.org/wiki/Overpass_API/Language_Guide
- [14] (2013, Apr.) OpenStreetMap wiki about the tagging system. [Online]. Available: http://wiki.openstreetmap.org/wiki/Map_Features
- [15] (2013, Apr.) Google Maps. <http://maps.google.com/>
- [16] (2013, Apr.) Wikipedia API. http://www.mediawiki.org/wiki/API:Main_page
- [17] (2013, Apr.) YouTube. <http://www.youtube.com/>